

FIVE  
AI

# Training neural networks using Tensorflow

Lars Mennen

Cambridge Wireless, 23rd November 2017



# Training neural networks using Tensorflow

- Two parts today:
  - General part about neural networks and how to train them
  - Training neural networks using Tensorflow
- Keep your laptops ready!
- Follow along with Python notebooks:

[https://github.com/larsmennen/intro\\_to\\_tensorflow](https://github.com/larsmennen/intro_to_tensorflow)

(most code adopted from Tensorflow tutorials - tensorflow.org)

# Part I:

# Training neural networks

# Supervised **learning**

- Today's focus: supervised learning

Given input pairs  
 $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$

where  $x_i$  comes from some input space  $X$  and  $y_i$  comes from some output space  $Y$ , we try to learn a function  $f$  such that:

$$f(x') = y'$$

for some unseen pair  $(x', y')$  (but from the same spaces!)

# Classifying **images**

- Basic example: image classification



Predict category for  
unseen image

Training data: 1.2M images + their categories  
(container ship, motor scooter, mushroom, ...)

2012 ImageNet classification  
challenge: only 16% top-5 error!

# Supervised learning **methods**

- Various algorithms that attempt to solve this problem:
  - Nearest neighbour
  - Decision tree learning
  - Support vector machines
  - **Neural Networks**
  - ...
- There are **many** approaches to supervised learning
- Neural networks are not always the answer



# Neural **networks**

- Inspired by biological **neurons**
- Main structure:
  - Relatively simple neurons that compute a function given some inputs
  - Structured in ordered layers, where the neurons in each layer have as input a weighted sum of outputs of neurons in the previous layers.



# Neural networks

- So, we want to learn  $f(x) = y$ . Usually  $x$  and  $y$  are represented as vectors.
- We feed  $x$  to the input layer.



For each neuron:


$$\text{activation} = g \left( \sum_{i=1}^n w_i a_i + b \right)$$



# Neural networks



- Popular choices for activation function  $g$ :
  - Sigmoid:

$$g(x) = \frac{1}{1 + e^{-x}}$$



- ReLU (rectified linear unit):

$$g(x) = \max(0, x)$$




$$\text{activation} = g \left( \sum_{i=1}^n w_i a_i + b \right)$$

Note that these are **non-linear**!

# Neural networks

- Given a neural network as on the right, an input  $x$  and a function  $g$  we can now compute the value of the node(s) in the output layer!
- We want this value to correspond to the label  $y$  in the pair  $(x,y)$ , as then the network is computing  $f(x) = y$ .




$$\text{activation} = g \left( \sum_{i=1}^n w_i a_i + b \right)$$



# Neural networks


- However, how do we know which:
  - Layer structure
  - Activation function  $g$
  - Values for weights for each layer
- we need to pick so that this network computes the function  $f$  that we want?



$$\text{activation} = g \left( \sum_{i=1}^n w_i a_i + b \right)$$

# Neural networks


- Unfortunately, don't have learning algorithm to find layer structure for hidden layers and/or activation function  $g$ .
- Found by reasoning, experimentation and building on previous research.
- In this talk we assume the structure and activation function are given.





# Neural networks

- We do have an algorithm to learn the weights:
  - Backpropagation
- The backpropagation algorithm, together with large amounts of data, powerful GPUs and *convolutional* neural networks (see later) is what makes modern NNs so popular and effective.



$$\text{activation} = g \left( \sum_{i=1}^n w_i a_i + b \right)$$

# Backpropagation **in a nutshell**

- We define a *loss function*:
  - Tells us "how far our prediction  $f(x) = y'$  is off" from the label  $y$
  - E.g. if we have a training example  $(x_i, y_i)$ , one possible loss function is:

$$L = \frac{1}{2} (y_i - f(x_i))^2$$

- We want to **minimise L**
- Parameter space is way too big to set derivatives equal to 0!



# Backpropagation **in a nutshell**

- So we do it iteratively.
- How do we know in which direction we have to change weights to minimize  $L$ ?
  - Look at the derivatives!
- Every single step in the neural network and the loss function are **differentiable**, which is key to the backpropagation algorithm.
- For every weight  $w_{ij}$  (from neuron  $i$  to neuron  $j$  in the next layer), we'd like to know:  $\frac{\partial L}{\partial w_{ij}}$  as then we know in which direction we should move  $w_{ij}$ .

# Backpropagation **in a nutshell**

- I'll skip the exact maths here (great explanation on <http://neuralnetworksanddeeplearning.com>), but main idea:
- If you know all the intermediate activations (i.e. all outputs of the activation function  $g$ ) for an input  $x_i$ , you can compute  $\frac{\partial L}{\partial w_{ij}}$  for all weights using the **chain rule for derivatives**.
- We can express gradients in a layer in terms of gradients of the next layer.
- So if we start at the last layer, we can **backpropagate** to find gradients in previous layers.



# Backpropagation **in a nutshell**

- Given these expressions, we can use an optimisation method, e.g. **gradient descent**, to adjust the weights in a way that will minimise the loss  $L$ .
- **Forward pass**: compute all activations for a given input  $x_i$ .
- **Backward pass**: compute gradients and change weights according to optimisation method

# Backpropagation **in a nutshell**

Main algorithm:

1. initialize all weights randomly
2. repeat until stopping criterion is met:
  - a. for  $(x_i, y_i)$  in dataset D:
    - i. **Forward pass:** compute  $f(x_i)$ , store intermediate activations, and compute  $L$
    - ii. **Backward pass:** compute gradients w.r.t.  $L$  and update weights according to optimisation method

Note that updating the weights changes  $f$ !


# Backpropagation **in a nutshell**

- In practice, do forward pass for multiple training examples at once (batching):
  - More efficient
  - Less noisy gradients
- Which stopping criterion to use?
  - Loss doesn't drop anymore
  - Better: look at loss on held-out validation set
    - Otherwise you might **overfit** on the particular training set, and hence **fail to generalise** to new examples.



# Using **convolutions**

- Fully connected layers don't scale well to images
- Convolutional layers:
  - Convolve learned weights with input
  - Weights shared along spatial dimensions
  - $k^2 \times c_i \times c_o$  weights for  $k \times k$  kernel from  $c_i$  input channels to  $c_o$  output channels



# Training **neural networks**

In summary:

- Given a supervised learning problem and a dataset  $D$ :
  - Define the structure of a neural net, an activation function and a loss function.
  - Learn the weights using the algorithm described before
  - You now have a function  $f: X \rightarrow Y$  which you can use to compute  $f(x)$  on unseen  $x$ .

# Part 2: Tensorflow



# Google **Tensorflow**

## What is Tensorflow?

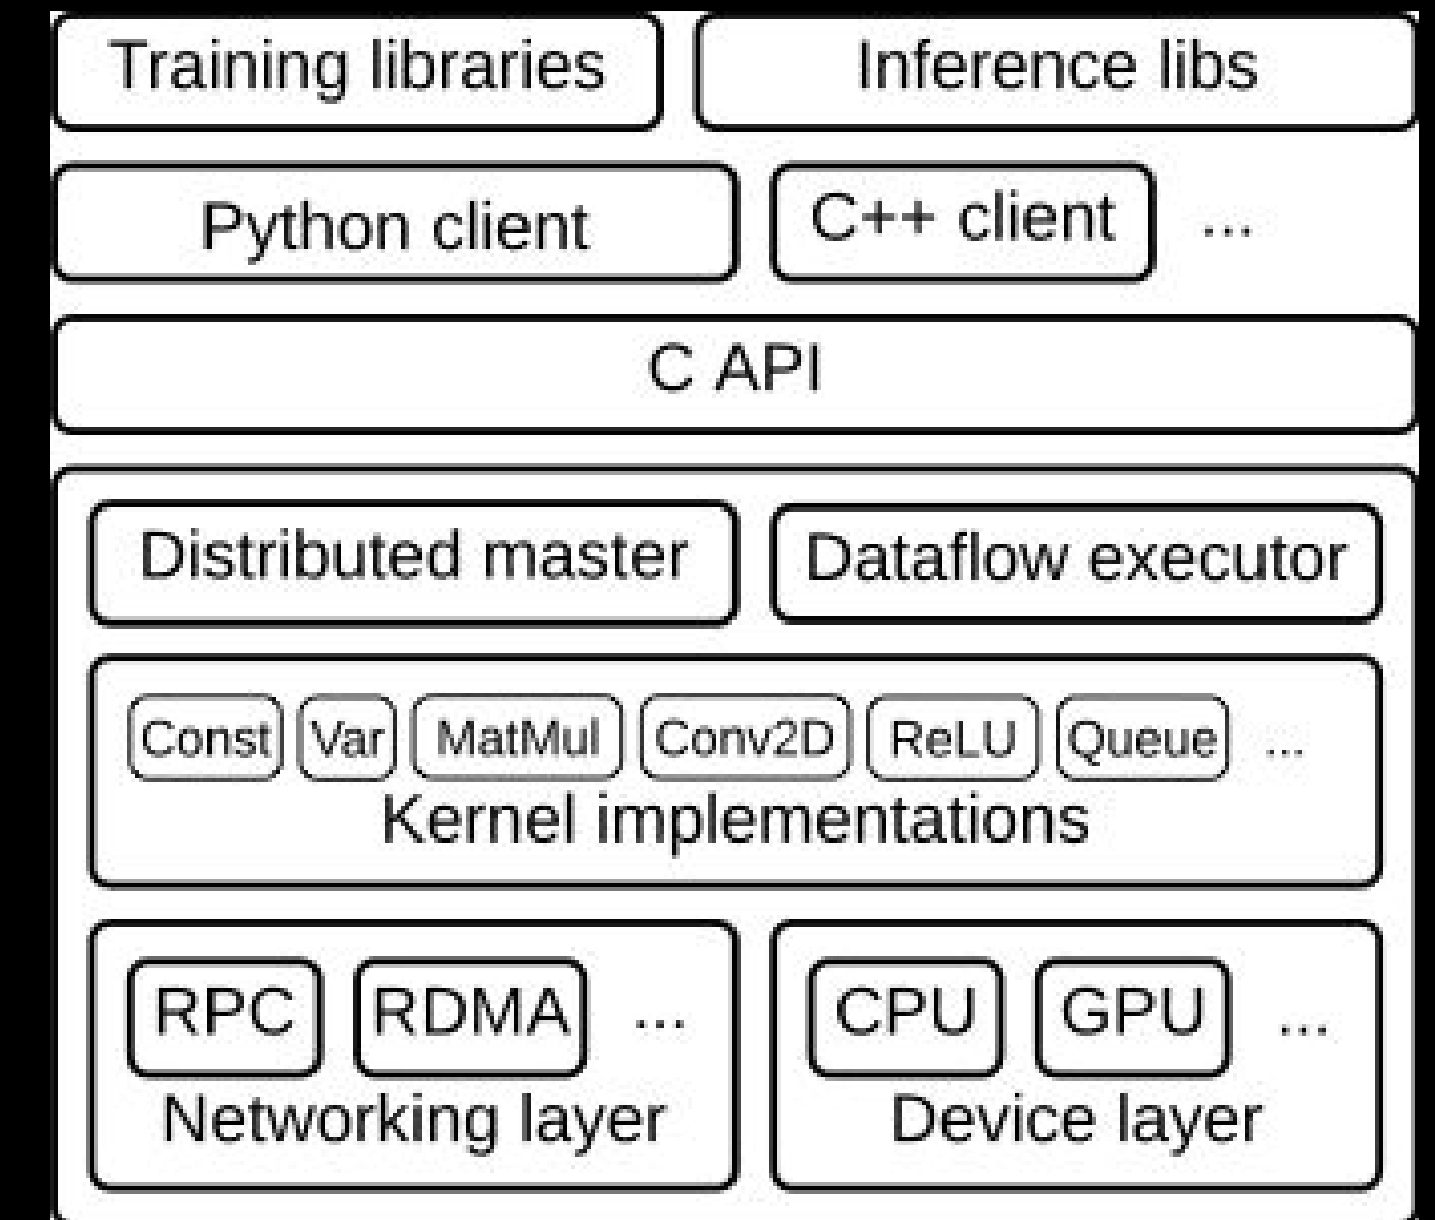
- *"TensorFlow™ is an open source software library for numerical computation using data flow graphs."*
- Probably the most popular open-source framework for training neural nets (but it's more general than that!)
- Large community, easy to use Python interface
- Used extensively in industry and research
- Development moves extremely fast!

# Tensorflow Overview

- Tensorflow allows you to define, train, evaluate and perform inference on neural networks.
- Lots of extra functionality:
  - Tensorboard - visualising neural networks and training
  - Serving - serving models in production
  - Training on HPC clusters
  - Preprocessing data
  - Quantization of neural networks
  - ...
- APIs for C++, Python, Java and Go

# Tensorflow **Architecture**

- Main implementations in C(++)
- Every operation can have a CPU and/or GPU implementation
- Most GPU code uses NVIDIA CUDA (proprietary)
  - CuDNN for common neural net operations
  - Efforts to get OpenCL support
- Relies heavily on **Eigen** and **Protobuf**



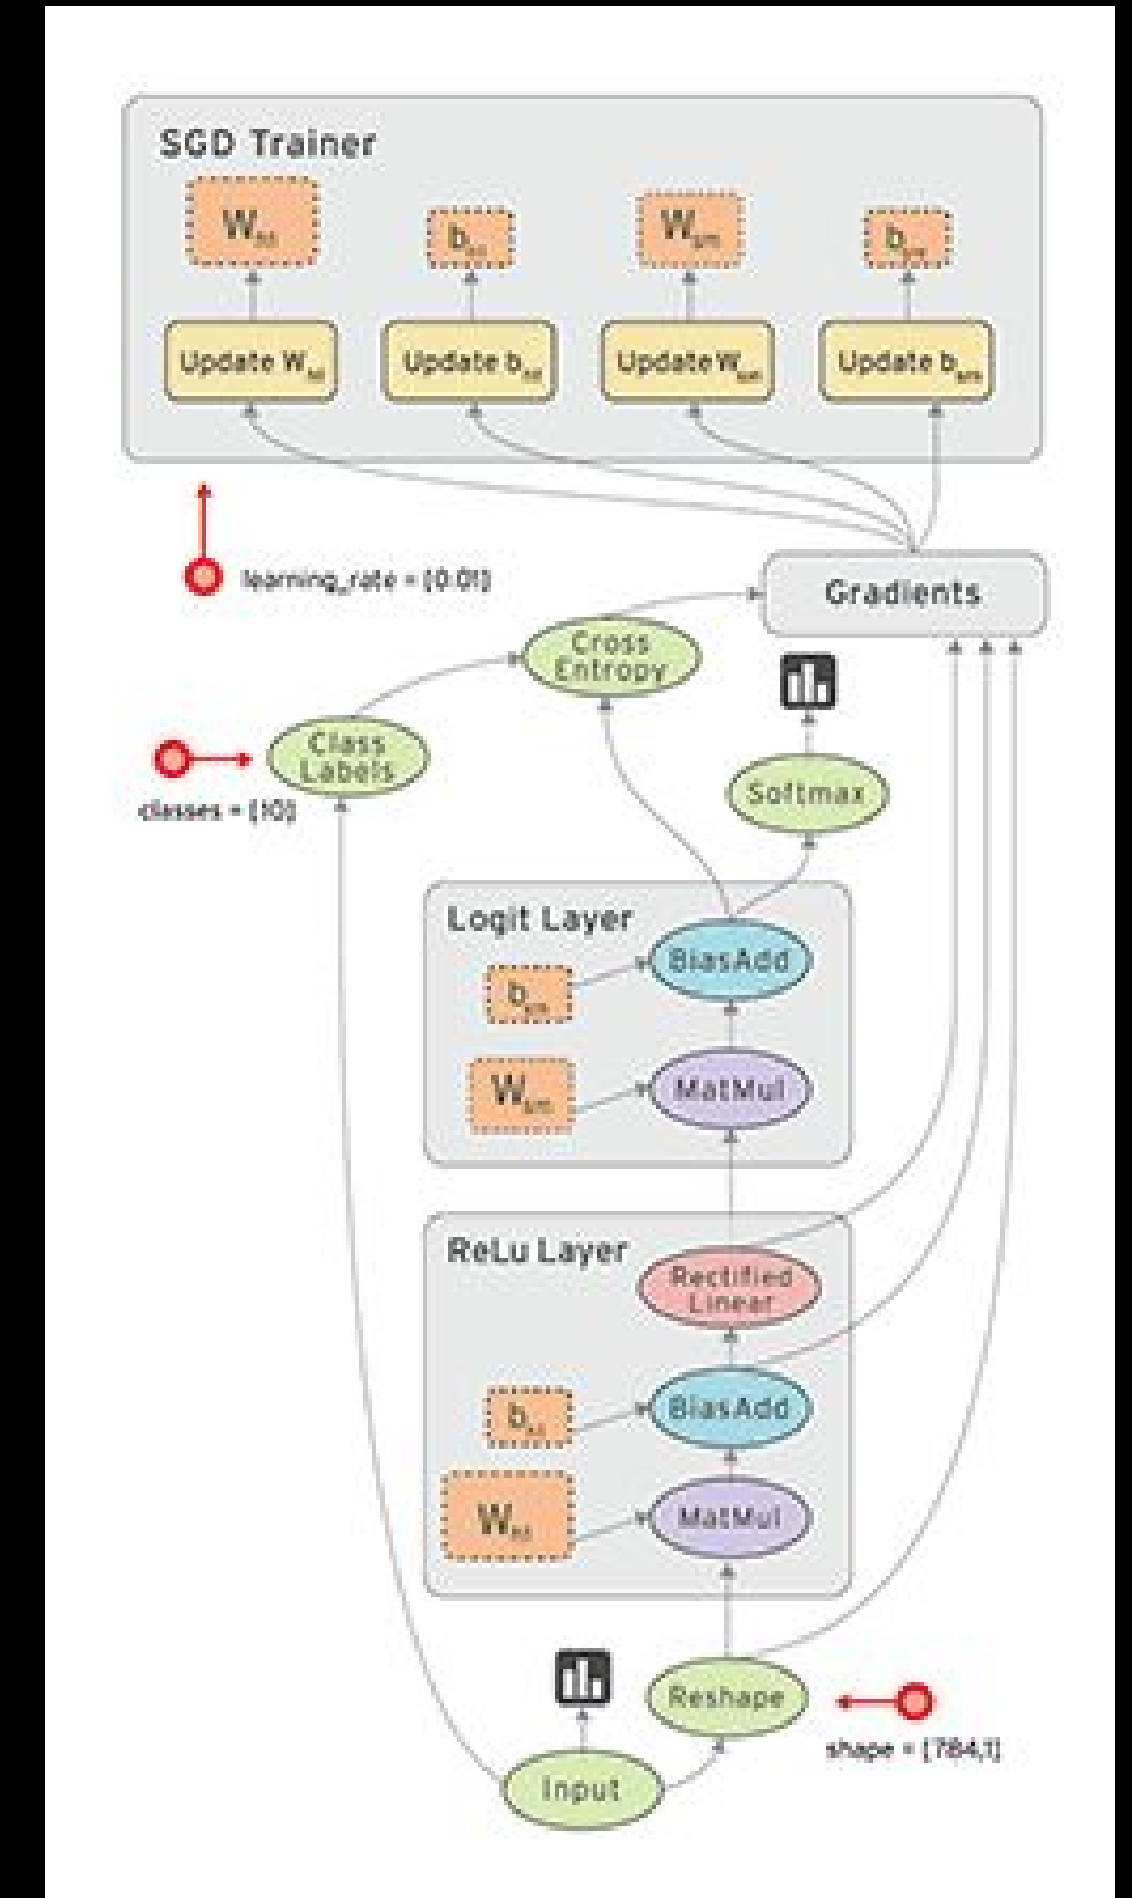


# Concepts: **Tensors**

- Computations in Tensorflow are done on **tensors**
- Generalisation of matrices to higher dimensions
- E.g. a tensor of rank 4 of dimensions (10,2,2,5) would have  $10 * 2 * 2 * 5 = 200$  elements
- Tensors have strong typing
- For input data, usually the first dimension is the **batch size**
  - E.g. feedforward pass for 4 images at once:  
(4, ...)

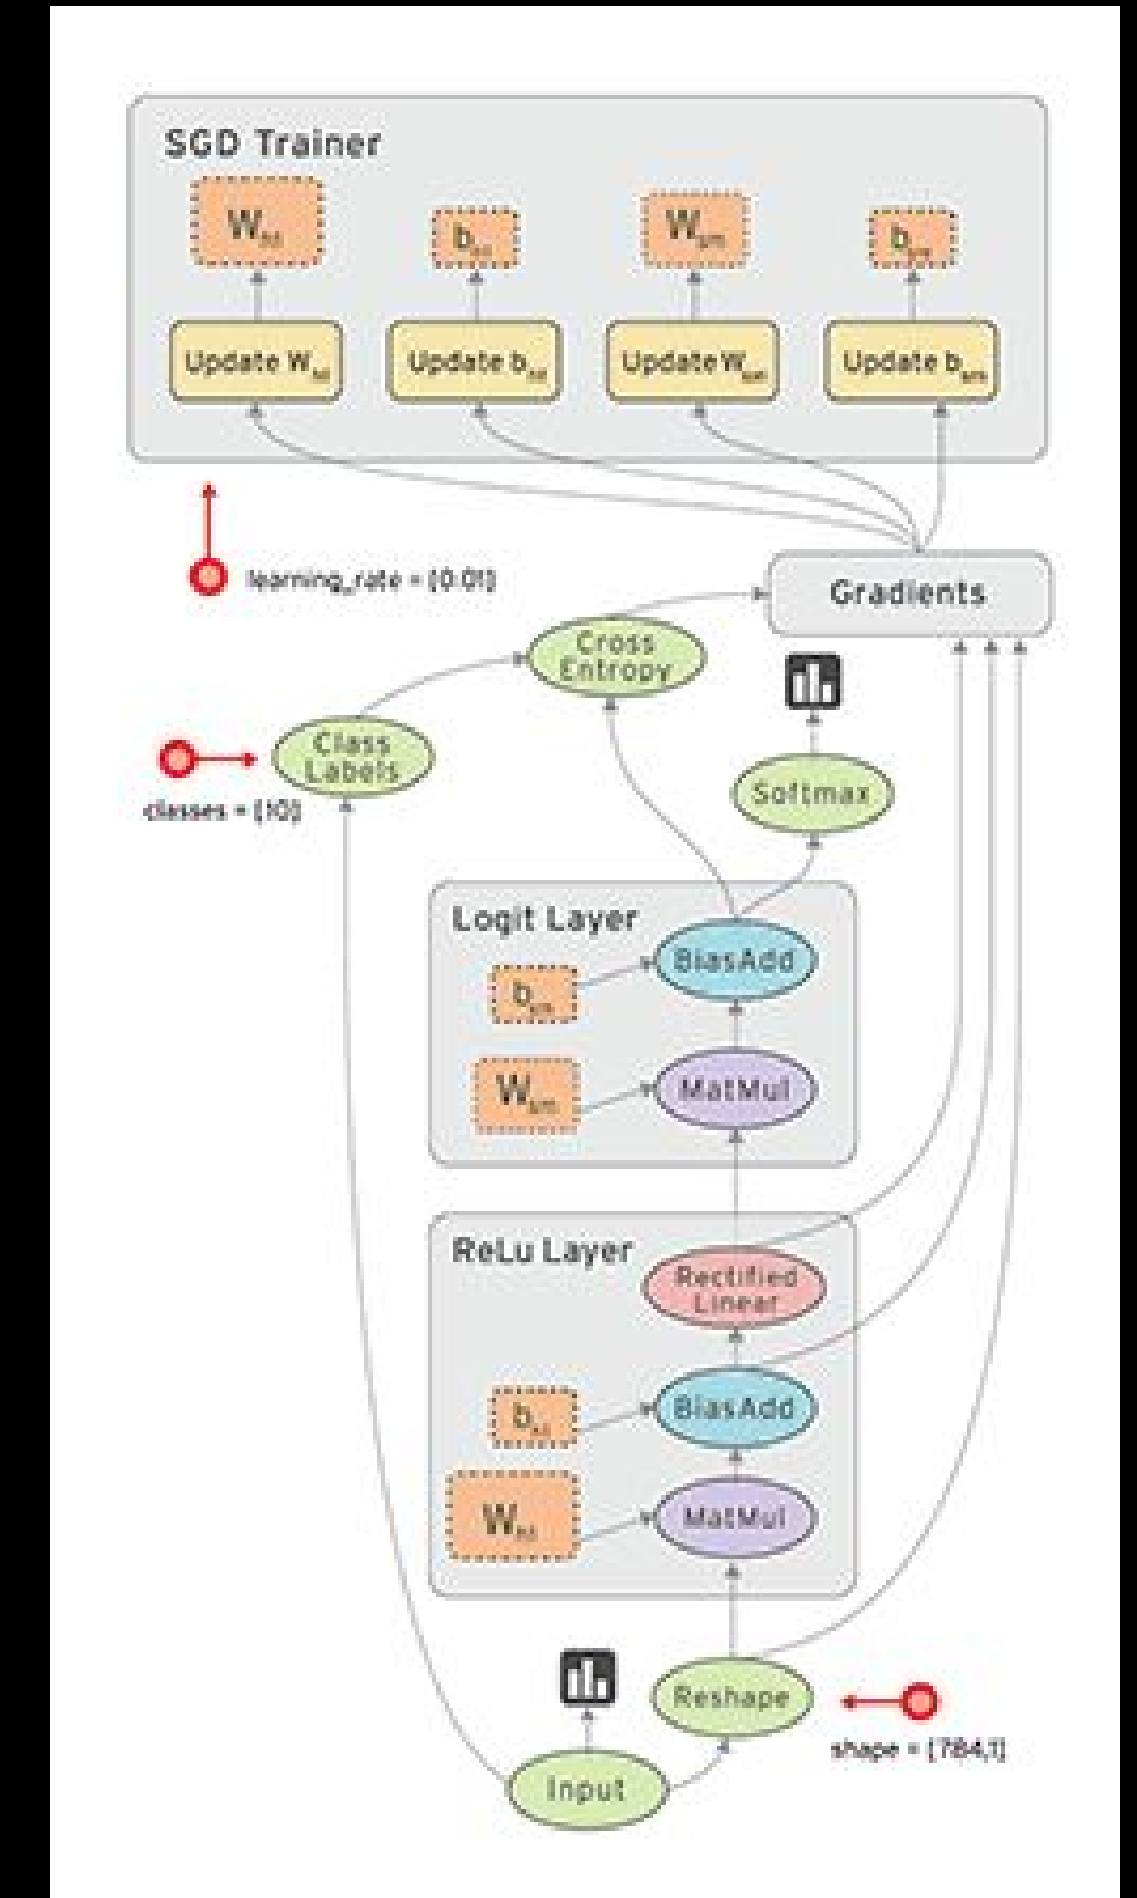
# Concepts: **Computation Graph**

- All computations in Tensorflow are represented in the computation graph
  - Neural network, optimiser, ...
- The majority of code you'll write in Python does not actually execute the network on data; it constructs the computation graph
- Graph consists of **Operations** whose inputs and outputs are **Tensors**.
- Input data is represented by placeholders



# Concepts: Operations and Kernels

- Operations run kernels
- Operations:
  - Metadata
  - Shape and type inference
    - Can work with partially defined shapes
  - Central registry
- Kernels:
  - Actual implementations on CPU or GPU
  - Can work for only certain types
  - Often Eigen for CPU kernels, NVIDIA CUDA/CuDNN for GPU kernels





# Concepts: **Operations and Kernels**

- Usually NN operations need **gradient operations**
- Tensorflow deduces which kernel to use and handles memory management for you
  - E.g. CPU-only operation after GPU-only operation
  - Possible to force placement on a specific CPU or GPU
- You can implement your own operations.
  - Python: as a combination of existing operations
  - C++: load at runtime as shared library

# Concepts: **Session**

- Represents the connection between the client (Python) and the C(++) runtime
- Provides access to the CPU and GPU device(s), which may be remote
- Allows to evaluate (parts of) the graph on data

Time to code!



# Installation



These instructions can be found on:

[https://github.com/larsmennen/intro\\_to\\_tensorflow](https://github.com/larsmennen/intro_to_tensorflow)

Vanilla Python or virtualenv (CPU only):

```
pip3 install tensorflow
```

Vanilla Python or virtualenv (GPU, CUDA and CuDNN present):

```
pip3 install tensorflow-gpu
```

Anaconda (NVIDIA GPU)

```
conda install tensorflow-gpu
```

Anaconda (no NVIDIA GPU)

```
conda install tensorflow
```

# Recognising Handwritten Digits

- We'll follow the Tensorflow tutorial on MNIST, but more in-depth
- Recognising handwritten digits
- **Classification problem**, 10 classes
- Data: pairs (x,y) where x is a 28x28 pixel image (which we'll flatten to a 784-element vector) of a handwritten digit and y is a 10-element one-hot vector representing the label
- 55k training, 5k validation, 10k test




# Network definition **in Tensorflow**

mnist.py and the inspecting\_mnist notebook



# Network definition in Tensorflow



Gives a total of:

$$784 * 128 + 128 * 32 + 32 * 10 = 104\,768$$

weights we need to train

Let's train our network!

# Network training in Tensorflow

training\_mnist notebook

# Higher level **interfaces**

- Using Tensorflow as we did today can get cumbersome
- There are higher level interfaces that make development easier and cleaner:
  - tf.slim
  - tf.estimator
  - Keras
- Keras provides a **clean, functional API** and only uses Tensorflow as a **backend** (can also use Microsoft CNTK or Theano)  
<https://keras.io/>



# More **resources**

- Explanation of (convolutional) neural networks:  
<http://neuralnetworksanddeeplearning.com>  
<http://cs231n.stanford.edu/>
- Tensorflow:  
<https://www.tensorflow.org/>
- OpenCL support for Tensorflow:  
<https://github.com/tensorflow/tensorflow/issues/22>



We're **hiring!**  
[five.ai/careers](https://five.ai/careers)

DevOps Engineer

Engineering Manager

Graduate Engineer

Internship

Machine Learning for Visual Scene Understanding

Machine Vision and Image Processing

Office Manager - Edinburgh

Research Engineer - Computer Vision and Deep Learning

Research Engineer - Robotics

Research Scientist

Research Scientist - Activity Understanding and Prediction

Research Scientist - Machine Learning for Driving Decisions

Research Scientist - Motion Prediction

Simulation Developer

Software Engineer - Machine Vision + Image Processing

Software Engineer - Platform

Questions?

# Operation **implementation**

Let's have a look under the hood. How is an operation actually implemented?

Simple example: `tf.sigmoid`  $g(x) = \frac{1}{1 + e^{-x}}$

`tf.sigmoid`

Aliases:

- `tf.nn.sigmoid`
- `tf.sigmoid`

```
sigmoid(  
    x,  
    name=None  
)
```

Defined in [tensorflow/python/ops/math\\_ops.py](#).

See the guide: [Neural Network > Activation Functions](#)

Computes sigmoid of `x` element-wise.

Specifically,  $y = 1 / (1 + \exp(-x))$ .



# Operation **implementation**

tensorflow/python/ops/math\_ops.py

```
def sigmoid(x, name=None):
    """Computes sigmoid of `x` element-wise.

    Specifically, `y = 1 / (1 + exp(-x))`.

    Args:
        x: A Tensor with type `float32`, `float64`, `int32`,
        `complex64`, `int64`,
        or `qint32`.
        name: A name for the operation (optional).

    Returns:
        A Tensor with the same type as `x` if `x.dtype !=
        qint32`
        otherwise the return type is `quint8`.

    @compatibility(numpy)
    Equivalent to np.scipy.special.expit
    @end_compatibility
    """
    with ops.name_scope(name, "Sigmoid", [x]) as name:
        x = ops.convert_to_tensor(x, name="x")
        return gen_math_ops._sigmoid(x, name=name)
```

Namespacing



If x is not already a tensor,  
convert it to a tensor



# Operation **implementation**

tensorflow/python/ops/gen\_math\_ops.py

```
def _sigmoid(x, name=None):  
    r"""Computes sigmoid of `x` element-wise.  
  
    Specifically,  $y = 1 / (1 + \exp(-x))$ .  
  
    Args:  
        x: A `Tensor`. Must be one of the following types:  
        `half`, `float32`, `float64`, `complex64`, `complex128`.  
        name: A name for the operation (optional).  
  
    Returns:  
        A `Tensor`. Has the same type as `x`.  
    """  
    result = op_def_lib.apply_op("Sigmoid", x=x, name=name)  
    return result
```

Just invokes the "Sigmoid"  
operation from the *operation  
registry* on *x*

Which will bring us to C++, so in Python only some  
conversions and checks!

# Operation **implementation**

tensorflow/core/ops/math\_ops.cc

```
REGISTER_OP("Sigmoid").UNARY_COMPLEX().Doc(R"doc(
Computes sigmoid of `x` element-wise.

Specifically, `y = 1 / (1 + exp(-x))`.
)doc");

REGISTER_OP("SigmoidGrad").UNARY_GRADIENT_COMPLEX().Doc(R"doc(
Computes the gradient of the sigmoid of `x` wrt its input.

Specifically, `grad = dy * y * (1 - y)`, where `y = sigmoid(x)`, and
`dy` is the corresponding input gradient.
)doc");
```

```
#define UNARY_COMPLEX() \
    Input("x: T") \
    .Output("y: T") \
    .Attr("T: {half, float, double, complex64, complex128}") \
    .SetShapeFn(shape_inference::UnchangedShape)

#define UNARY_GRADIENT_COMPLEX() \
    Input("x: T") \
    .Input("y: T") \
    .Output("z: T") \
    .Attr("T: {half, float, double, complex64, complex128}") \
    .SetShapeFn(shape_inference::UnchangedShape)
```

This registers the **Operation**, but no **kernels** yet. So this is just "metadata".  
Note both Sigmoid and SigmoidGrad.

# Operation **implementation**

tensorflow/core/kernels/cwise\_op\_sigmoid.cc

```
REGISTER5(UnaryOp, CPU, "Sigmoid", functor::sigmoid, float, Eigen::half, double,
          complex64, complex128);
#ifdef GOOGLE_CUDA
REGISTER3(UnaryOp, GPU, "Sigmoid", functor::sigmoid, float, Eigen::half,
          double);
#endif
```

```
REGISTER5(SimpleBinaryOp, CPU, "SigmoidGrad", functor::sigmoid_grad, float,
          Eigen::half, double, complex64, complex128);
#ifdef GOOGLE_CUDA
REGISTER3(SimpleBinaryOp, GPU, "SigmoidGrad", functor::sigmoid_grad, float,
          Eigen::half, double);
#endif
```

```
#define REGISTER(OP, D, N, F, T) \
    REGISTER_KERNEL_BUILDER(Name(N).Device(DEVICE_##D).TypeConstraint<T>("T"), \
                             OP<D##Device, F<T>>);
```

This registers the kernels. Separate for CPU and GPU, may have different supported features.